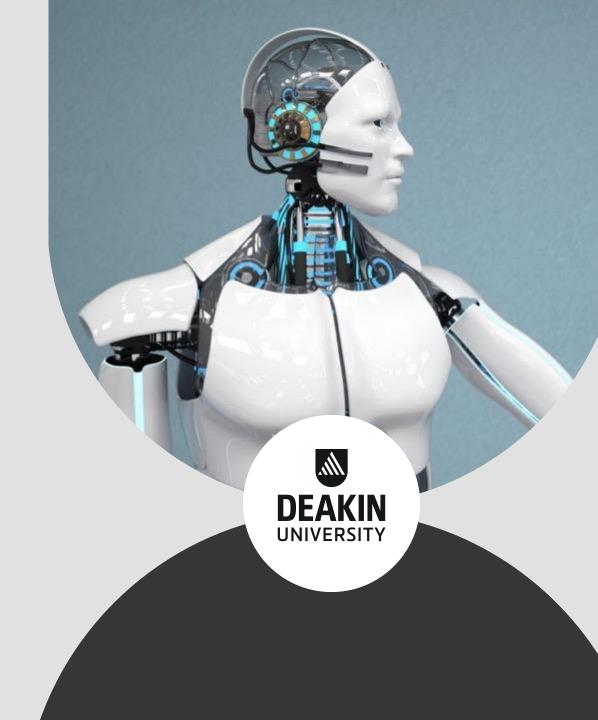
SIT330-770: Natural Language Processing

Week 5 - Vector Embeddings and Sequence Labeling

Dr. Mohamed Reda Bouadjenek

School of Information Technology, Faculty of Sci Eng & Built Env

reda.bouadjenek@deakin.edu.au

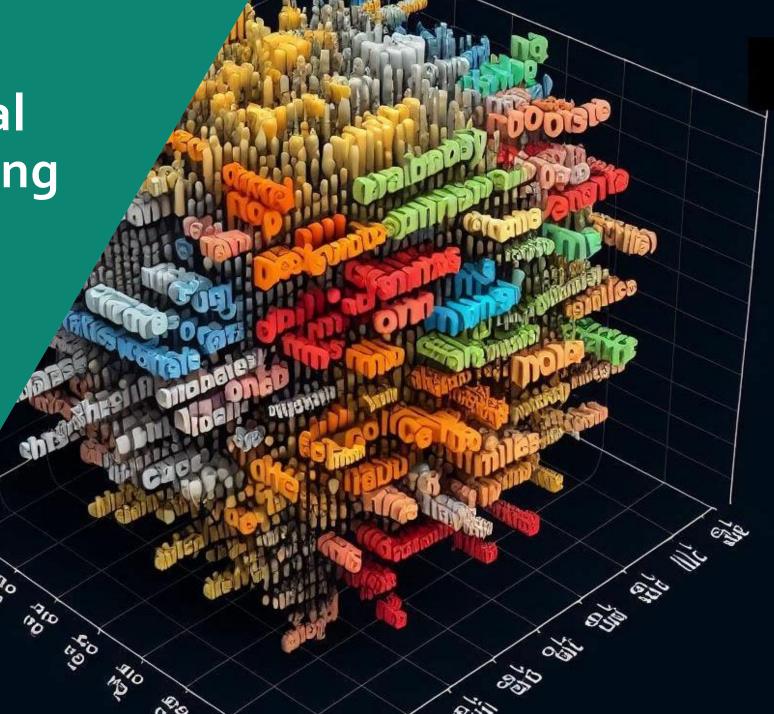


SIT330-770: Natural Language Processing

Week 5.1 - Word Meaning

Dr. Mohamed Reda Bouadjenek

School of Information Technology, Faculty of Sci Eng & Built Env



What do words mean?



- N-gram or text classification methods we've seen so far
 - Words are just strings (or indices w_i in a vocabulary list)
 - That's not very satisfactory!
- Introductory logic classes:
 - The meaning of "dog" is DOG; cat is CAT $\forall x DOG(x) \longrightarrow MAMMAL(x)$
- Old linguistics joke by Barbara Partee in 1967:
 - O: What's the meaning of life?
 - O A: LIFE
- That seems hardly better!

Desiderata



- What should a theory of word meaning do for us?
- Let's look at some desiderata
- From lexical semantics, the linguistic study of word meaning

Lemmas and senses





mouse (N)

sense

- 1. any of numerous small rodents...
- 2. a hand-operated device that controls

a cursor...

Modified from the online thesaurus WordNet

A sense or "concept" is the meaning component of a word Lemmas can be polysemous (have multiple senses)

Relations between senses: Synonymy



- Synonyms have the same meaning in some or all contexts.
 - o filbert / hazelnut
 - o couch / sofa
 - obig / large
 - o automobile / car
 - ovomit / throw up
 - o water / H₂o

Relations between senses: Synonymy



 Note that there are probably no examples of perfect synonymy.

- Even if many aspects of meaning are identical
- Still may differ based on politeness, slang, register, genre, etc.

Relation: Synonymy?



water/H₂o

"H₂o" in a surfing guide?

big/large

my big sister != my large sister

The Linguistic Principle of Contrast



• Difference in form \rightarrow difference in meaning

Abbé Gabriel Girard 1718



Re: "exact" synonyms

je ne crois pas qu'il y ait demot synonime dans aucune. Langue le le dis par con-

[I do not believe that there is a synonymous word in any language]

LA JUSTESSE

DE LA

LANGUE FRANÇOISE.

ου

LES DIFFERENTES SIGNIFICATIONS

DES MOTS QUIPASSENT

POUR

SYNONIMES

Par M. l'Abbé GIRARD C. D. M. D. D. B.



A PARIS,

Chez LAURENT D'HOURY, Imprimeur-L'braire, au bas de la rue de la Harpe, visà vis la rue S. Severin, au Saint Esprit.

M. DCC. XVIII.

Avec Approbation & Privilega dis Roy.

Relation: Similarity



Words with similar meanings. Not synonyms, but sharing some element of meaning

car, bicycle

cow, horse

Ask humans how similar 2 words are



word1	word2	similarity
vanish	disappear	9.8
behave	obey	7.3
belief	impression	5.95
muscle	bone	3.65
modest	flexible	0.98
hole	agreement	0.3

Relation: Word relatedness



- Also called "word association"
- Words can be related in any way, perhaps via a semantic frame or field

```
ocoffee, tea: similar
```

ocoffee, cup: related, not similar

Semantic field



- Words that
 - o cover a particular semantic domain
 - o bear structured relations with each other.

hospitals

surgeon, scalpel, nurse, anaesthetic, hospital

restaurants

waiter, menu, plate, food, menu, chef

houses

door, roof, kitchen, family, bed

Relation: Antonymy



- Senses that are opposites with respect to only one feature of meaning
- Otherwise, they are very similar!

```
dark/light short/long fast/slowrise/fall
hot/cold up/down in/out
```

- More formally: antonyms can
 - o define a binary opposition or be at opposite ends of a scale
 - o long/short, fast/slow
 - Be reversives:
 - o rise/fall, up/down

Connotation (sentiment)



- Words have affective meanings
 - Positive connotations (happy)
 - Negative connotations (sad)
- Connotations can be subtle:
 - Positive connotation: copy, replica, reproduction
 - Negative connotation: fake, knockoff, forgery
- Evaluation (sentiment!)
 - Positive evaluation (great, love)
 - Negative evaluation (*terrible*, *hαte*)

Connotation



Osgood et al. (1957)

- Words seem to vary along 3 affective dimensions:
 - o valence: the pleasantness of the stimulus
 - o arousal: the intensity of emotion provoked by the stimulus
 - o dominance: the degree of control exerted by the stimulus

	Word	Score	Word	Score
Valence	love	1.000	toxic	0.008
	happy	1.000	nightmare	0.005
Arousal	elated	0.960	mellow	0.069
	frenzy	0.965	napping	0.046
Dominance	powerful	0.991	weak	0.045
	leadership	0.983	empty	0.081

So far



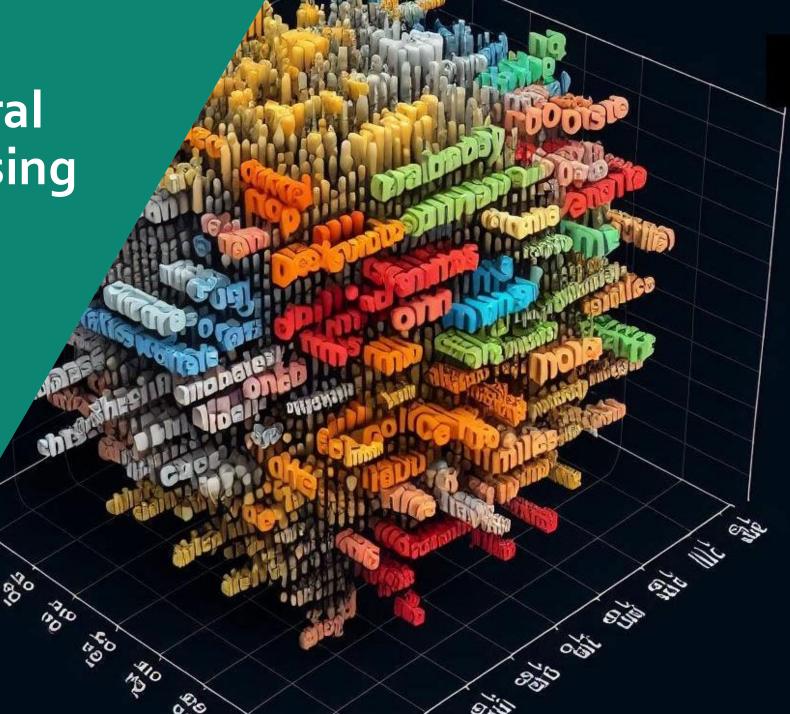
- Concepts or word senses
 - Have a complex many-to-many association with words (homonymy, multiple senses)
- Have relations with each other
 - Synonymy
 - Antonymy
 - Similarity
 - Relatedness
 - Connotation

SIT330-770: Natural Language Processing

Week 5.2 - Vector Semantics

Dr. Mohamed Reda Bouadjenek

School of Information Technology, Faculty of Sci Eng & Built Env



Computational models of word meaning



- Can we build a theory of how to represent word meaning, that accounts for at least some of the desiderata?
- We'll introduce vector semantics
 - The standard model in language processing!
 - Handles many of our goals!

Ludwig Wittgenstein



•PI #43:

"The meaning of a word is its use in the language"

Let's define words by their usages



- One way to define "usage":
- words are defined by their environments (the words around them)

- Zellig Harris (1954):
- If A and B have almost identical environments we say that they are synonyms.

What does recent English borrowing ongchoi mean?



- Suppose you see these sentences:
 - Ong choi is delicious sautéed with garlic.
 - Ong choi is superb over rice
 - Ong choi leaves with salty sauces
- And you've also seen these:
 - ...spinach sautéed with garlic over rice
 - Chard stems and leaves are delicious
 - Collard greens and other salty leafy greens
- Conclusion:
 - Ongchoi is a leafy green like spinach, chard, or collard greens
 - We could conclude this based on words like "leaves" and "delicious" and "sauteed"

Ongchoi: *Ipomoea aquatica "Water Spinach"*



空心菜 kangkong rau muống



Idea 1: Defining meaning by linguistic distribution



• Let's define the meaning of a word by its distribution in language use, meaning its neighboring words or grammatical environments.

Idea 2: Meaning as a point in space (Osgood et al. 1957)



3 affective dimensions for a word

valence: pleasantness

o **arousal**: intensity of emotion

o **dominance**: the degree of control exerted

	Word	Score	Word	Score
Valence	love	1.000	toxic	0.008
	happy	1.000	nightmare	0.005
Arousal	elated	0.960	mellow	0.069
	frenzy	0.965	napping	0.046
Dominance	powerful	0.991	weak	0.045
	leadership	0.983	empty	0.081

NRC VAD Lexicon (Mohammad 2018)

0

Hence the connotation of a word is a vector in 3-space



Idea 1: Defining meaning by linguistic distribution

Idea 2: Meaning as a point in multidimensional space

Defining meaning as a point in space based on distribution



- Each word = a vector (not just "good" or "w₄₅")
- Similar words are "nearby in semantic space"
- We build this space automatically by seeing which words are nearby in text

```
not good
                                                            bad
                                                  dislike
to
       by
                                                                 worst
                   's
                                                 incredibly bad
that
        now
                      are
                                                                   worse
                 vou
 than
         with
                  is
                                          incredibly good
                             very good
                      amazing
                                         fantastic
                                                  wonderful
                  terrific
                                      nice
                                     good
```

We define meaning of a word as a vector



- Called an "embedding" because it's embedded into a space (see textbook)
- The standard way to represent meaning in NLP
- Every modern NLP algorithm uses embeddings as the representation of word meaning
- Fine-grained model of meaning for similarity

Intuition: why vectors?



- Consider sentiment analysis:
 - With words, a feature is a word identity
 - Feature 5: 'The previous word was "terrible"
 - o requires **exact same word** to be in training and test

With embeddings:

- Feature is a word vector
- The previous word was vector [35,22,17...]
- Now in the test set we might see a similar vector [34,21,14]
- We can generalize to similar but unseen words!!!

We'll discuss 2 kinds of embeddings



tf-idf

- Information Retrieval workhorse!
- A common baseline model
- Sparse vectors
- Words are represented by (a simple function of) the counts of nearby words

Word2vec

- Dense vectors
- Representation is created by training a classifier to predict whether a word is likely to appear nearby
- Later we'll discuss extensions called contextual embeddings

From now on: Computing with meaning representations instead of string representations



荃者所以在鱼,得鱼而忘荃 Nets are for fish;

言者所以在意,得意而忘言 Words are for meaning;

Once you get the meaning, you can forget the words 庄子(Zhuangzi), Chapter 26

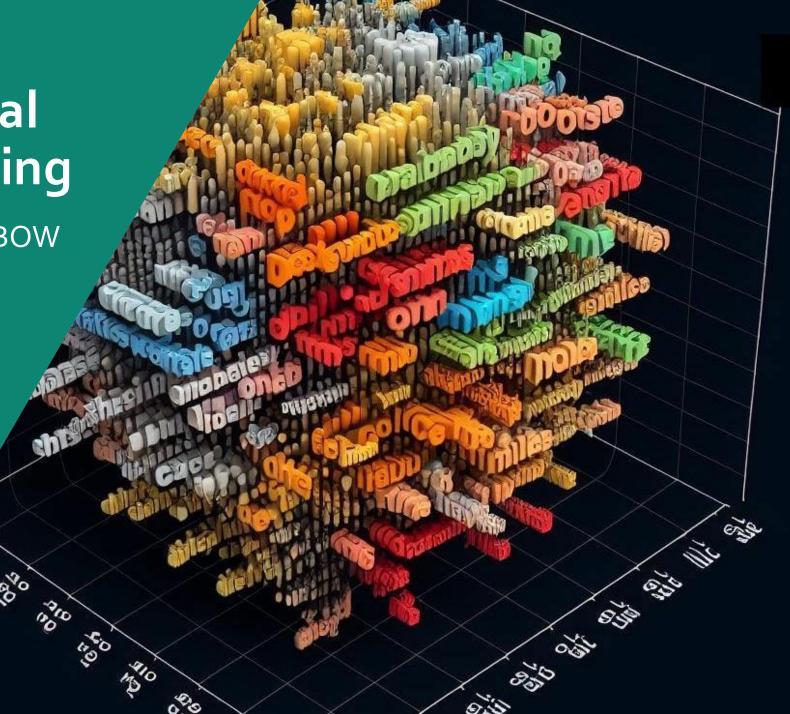
Once you get the fish, you can forget the net.



Week 5.3 – Words and Vectors: BOW

Dr. Mohamed Reda Bouadjenek

School of Information Technology, Faculty of Sci Eng & Built Env



Bag of Words



- A document is represented as vector of words.
 - One dimension per word.
 - Vector size is the vocabulary size, e.g., English may contain 100k words.
 - o Different weighting schemas can be used, e.g., tf, log(tf), tf-idf, Boolean, etc.
 - Sparse vector, e.g., almost all values are zeros.

	As You Like It	Twelfth Night	Julius Caesar	Henry V
battle		0	7	13
good	14	80	62	89
fool	36	58	1	4
wit	20	15	2	3

Order matters for NLP tasks!



- Assumes independence between words:
 - The sentences "John likes Mary" has the same representation as "Mary likes John" –
 even though the semantic is different).
- May work well for Information Retrieval tasks, but not for NLP tasks!
 - Sentiment analysis:

"Ah no, there are good movies on Netflix!" vs. "Ah, there are no good movies on Netflix!"

Computing word similarity: Dot product and cosine



The dot product between two vectors is a scalar:

$$dot product(\mathbf{v}, \mathbf{w}) = \mathbf{v} \cdot \mathbf{w} = \sum_{i=1}^{N} v_i w_i = v_1 w_1 + v_2 w_2 + \dots + v_N w_N$$

- The dot product tends to be high when the two vectors have large values in the same dimensions
- Dot product can thus be a useful similarity metric between vectors

Problem with raw dot-product



- Dot product favors long vectors
- Dot product is higher if a vector is longer (has higher values in many dimension)
- Vector length:

$$|\mathbf{v}| = \sqrt{\sum_{i=1}^{N} v_i^2}$$

- Frequent words (of, the, you) have long vectors (since they occur many times with other words).
- So dot product overly favors frequent words

Alternative: cosine for computing word similarity



$$cosine(\vec{v}, \vec{w}) = \frac{\vec{v} \cdot \vec{w}}{|\vec{v}||\vec{w}|} = \frac{\sum_{i=1}^{N} v_i w_i}{\sqrt{\sum_{i=1}^{N} v_i^2 \sqrt{\sum_{i=1}^{N} w_i^2}}}$$

Based on the definition of the dot product between two vectors a and b

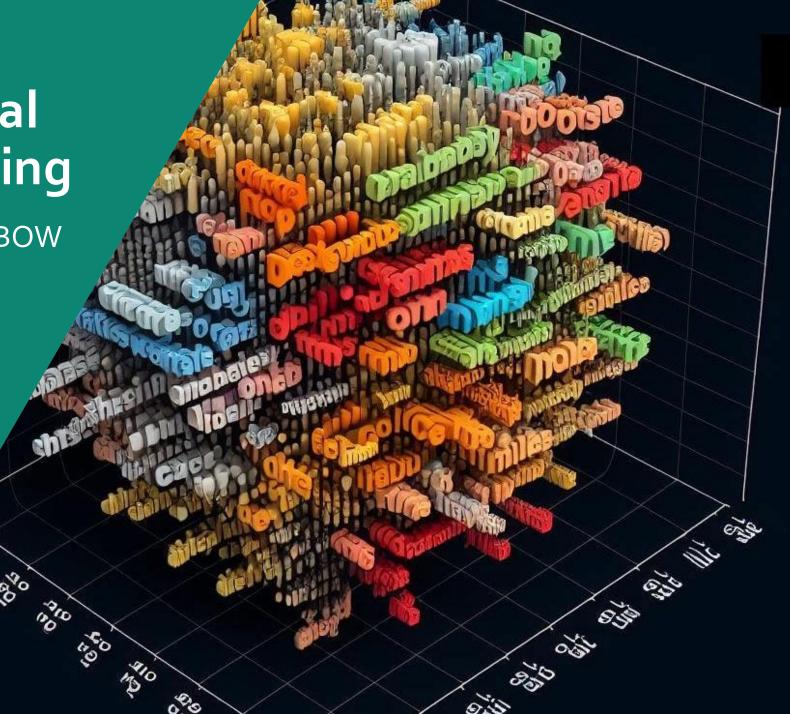
$$\frac{\mathbf{a} \cdot \mathbf{b}}{\frac{\mathbf{a} \cdot \mathbf{b}}{|\mathbf{a}||\mathbf{b}|}} = \cos \theta$$



Week 5.4 – Words and Vectors: BOW

Dr. Mohamed Reda Bouadjenek

School of Information Technology, Faculty of Sci Eng & Built Env



Bag of Words



- A document is represented as vector of words.
 - One dimension per word.
 - Vector size is the vocabulary size, e.g., English may contain 100k words.
 - o Different weighting schemas can be used, e.g., tf, log(tf), tf-idf, Boolean, etc.
 - Sparse vector, e.g., almost all values are zeros.

	As You Like It	Twelfth Night	Julius Caesar	Henry V
battle		0	7	13
good	14	80	62	89
fool	36	58	1	4
wit	20	15	2	3

Order matters for NLP tasks!



- Assumes independence between words:
 - The sentences "John likes Mary" has the same representation as "Mary likes John" –
 even though the semantic is different).
- May work well for Information Retrieval tasks, but not for NLP tasks!
 - Sentiment analysis:

"Ah no, there are good movies on Netflix!" vs. "Ah, there are no good movies on Netflix!"

Computing word similarity: Dot product and cosine



The dot product between two vectors is a scalar:

$$dot product(\mathbf{v}, \mathbf{w}) = \mathbf{v} \cdot \mathbf{w} = \sum_{i=1}^{N} v_i w_i = v_1 w_1 + v_2 w_2 + \dots + v_N w_N$$

- The dot product tends to be high when the two vectors have large values in the same dimensions
- Dot product can thus be a useful similarity metric between vectors

Problem with raw dot-product



- Dot product favors long vectors
- Dot product is higher if a vector is longer (has higher values in many dimension)
- Vector length:

$$|\mathbf{v}| = \sqrt{\sum_{i=1}^{N} v_i^2}$$

- Frequent words (of, the, you) have long vectors (since they occur many times with other words).
- So dot product overly favors frequent words

Alternative: cosine for computing word similarity



$$cosine(\vec{v}, \vec{w}) = \frac{\vec{v} \cdot \vec{w}}{|\vec{v}||\vec{w}|} = \frac{\sum_{i=1}^{N} v_i w_i}{\sqrt{\sum_{i=1}^{N} v_i^2 \sqrt{\sum_{i=1}^{N} w_i^2}}}$$

Based on the definition of the dot product between two vectors a and b

$$\frac{\mathbf{a} \cdot \mathbf{b}}{|\mathbf{a}||\mathbf{b}|} = \cos \theta$$

Sparse versus dense vectors



- tf-idf (or PMI) vectors are
 - olong (length |V|= 20,000 to 50,000)
 - osparse (most elements are zero)
- Alternative: learn vectors which are
 - **short** (length 50-1000)
 - odense (most elements are non-zero)

Sparse versus dense vectors



Why dense vectors?

- Short vectors may be easier to use as features in machine learning (fewer weights to tune)
- Dense vectors may generalize better than explicit counts
- Dense vectors may do better at capturing synonymy:
 - ocar and automobile are synonyms; but are distinct dimensions
 - •a word with *cαr* as a neighbor and a word with *αutomobile* as a neighbor should be similar, but aren't
- In practice, they work better

Common methods for getting short dense vectors



- "Neural Language Model"-inspired models
 - Word2vec (skipgram, CBOW), GloVe
- Singular Value Decomposition (SVD)
 - A special case of this is called LSA Latent Semantic Analysis
- Alternative to these "static embeddings":
 - Contextual Embeddings (ELMo, BERT)
 - Compute distinct embeddings for a word in its context
 - Separate embeddings for each token of a word

Simple static embeddings you can download!



- Word2vec (Mikolov et al)
- https://code.google.com/archive/p/word2vec/

- GloVe (Pennington, Socher, Manning)
- http://nlp.stanford.edu/projects/glove/

Word2vec



- Popular embedding method
- Very fast to train
- Code available on the web
- Idea: predict rather than count
- Word2vec provides various options. We'll do:
- skip-gram with negative sampling (SGNS)

Word₂Vec



- Instead of counting how often each word w occurs near "αpricot"
 - Train a classifier on a binary prediction task:
 - Is w likely to show up near "apricot"?
- We don't actually care about this task
 - But we'll take the learned classifier weights as the word embeddings
- Big idea: self-supervision:
 - A word c that occurs near apricot in the corpus cats as the gold "correct answer" for supervised learning
 - No need for human labels
 - O Bengio et al. (2003); Collobert et al. (2011)

Approach: predict if candidate word c is a "neighbor"



- Treat the target word t and a neighboring context word c as positive examples.
- 2. Randomly sample other words in the lexicon to get negative examples
- 3. Use logistic regression to train a classifier to distinguish those two cases
- 4. Use the learned weights as the embeddings

Skip-Gram Training Data



(assuming a +/- 2 word window)

Goal: train a classifier that is given a candidate (word, context) pair

```
(apricot, jam)(apricot, aardvark)
```

• • •

- And assigns each pair a probability:
 - $\circ P(+|w,c)$
 - P(-|w, c) = 1 P(+|w, c)

Similarity is computed from dot product



- Remember: two vectors are similar if they have a high dot product
 - Cosine is just a normalized dot product
- So:
 - \circ Similarity(w,c) \propto w · c
- We'll need to normalize to get a probability
 - (cosine isn't a probability either)

Turning dot products into probabilities



- $Sim(w,c) \approx w \cdot c$
- To turn this into a probability
- We'll use the sigmoid from logistic regression:

$$P(+|w,c) = \sigma(c \cdot w) = \frac{1}{1 + \exp(-c \cdot w)}$$

$$P(-|w,c) = 1 - P(+|w,c)$$

$$= \sigma(-c \cdot w) = \frac{1}{1 + \exp(c \cdot w)}$$

How Skip-Gram Classifier computes P(+|w, c)



$$P(+|w,c) = \sigma(c \cdot w) = \frac{1}{1 + \exp(-c \cdot w)}$$

- This is for one context word, but we have lots of context words.
- We'll assume independence and just multiply them:

$$P(+|w,c_{1:L}) = \prod_{i=1}^{L} \sigma(c_i \cdot w)$$

$$\log P(+|w,c_{1:L}) = \sum_{i=1}^{L} \log \sigma(c_i \cdot w)$$

Skip-gram classifier: summary



- A probabilistic classifier, given
 - a test target word w
 - its context window of L words $c_{1:L}$
- Estimates probability that w occurs in this window based on similarity of w (embeddings) to $C_{1:L}$ (embeddings).

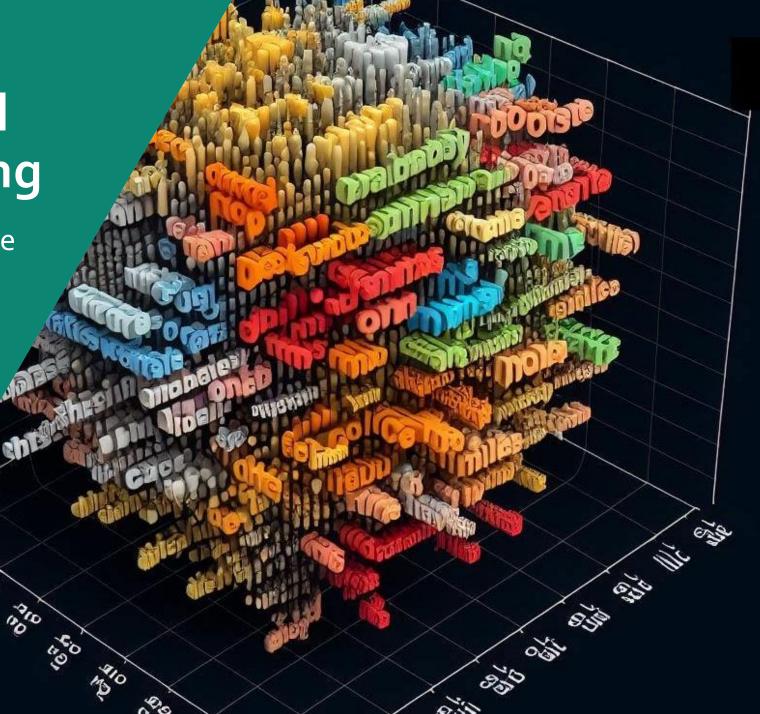
To compute this, we just need embeddings for all the words.



Week 5.5 — Word2vec: Learning the embeddings

Dr. Mohamed Reda Bouadjenek

School of Information Technology, Faculty of Sci Eng & Built Env



Skip-Gram Training data



...lemon, a [tablespoon of apricot jam, a] pinch...

C1

c2 [target] c3 c4

positive examples +

t

apricot tablespoon apricot of apricot jam apricot a

Skip-Gram Training data



...lemon, a [tablespoon of apricot jam, a] pinch...

C1

c2 [target] c3 c4

positive examples +

t c

apricot tablespoon apricot of apricot jam apricot a For each positive example we'll grab k negative examples, sampling by frequency

Skip-Gram Training data



...lemon, a [tablespoon of apricot jam, a] pinch...

C1

c2 [target] c3 c4

positive examples +

t

apricot tablespoon apricot of apricot jam apricot a

negative examples -

t	c	t	C
apricot	aardvark	apricot	seven
apricot	my	apricot	forever
apricot	where	apricot	dear
apricot	coaxial	apricot	if

Word2vec: how to learn vectors



- Given the set of positive and negative training instances, and an initial set of embedding vectors
- The goal of learning is to adjust those word vectors such that we:
 - Maximize the similarity of the target word, context word pairs (w, c_{pos}) drawn from the positive data
 - \circ Minimize the similarity of the (w, c_{neq}) pairs drawn from the negative data.

Loss function for one w with c_{pos} , c_{neg1} ...c_{negk}



 Maximize the similarity of the target with the actual context words, and minimize the similarity of the target with the k negative sampled non-neighbor words.

$$L_{CE} = -\log \left[P(+|w, c_{pos}) \prod_{i=1}^{k} P(-|w, c_{neg_i}) \right]$$

$$= -\left[\log P(+|w, c_{pos}) + \sum_{i=1}^{k} \log P(-|w, c_{neg_i}) \right]$$

$$= -\left[\log P(+|w, c_{pos}) + \sum_{i=1}^{k} \log \left(1 - P(+|w, c_{neg_i}) \right) \right]$$

$$= -\left[\log \sigma(c_{pos} \cdot w) + \sum_{i=1}^{k} \log \sigma(-c_{neg_i} \cdot w) \right]$$

Learning the classifier

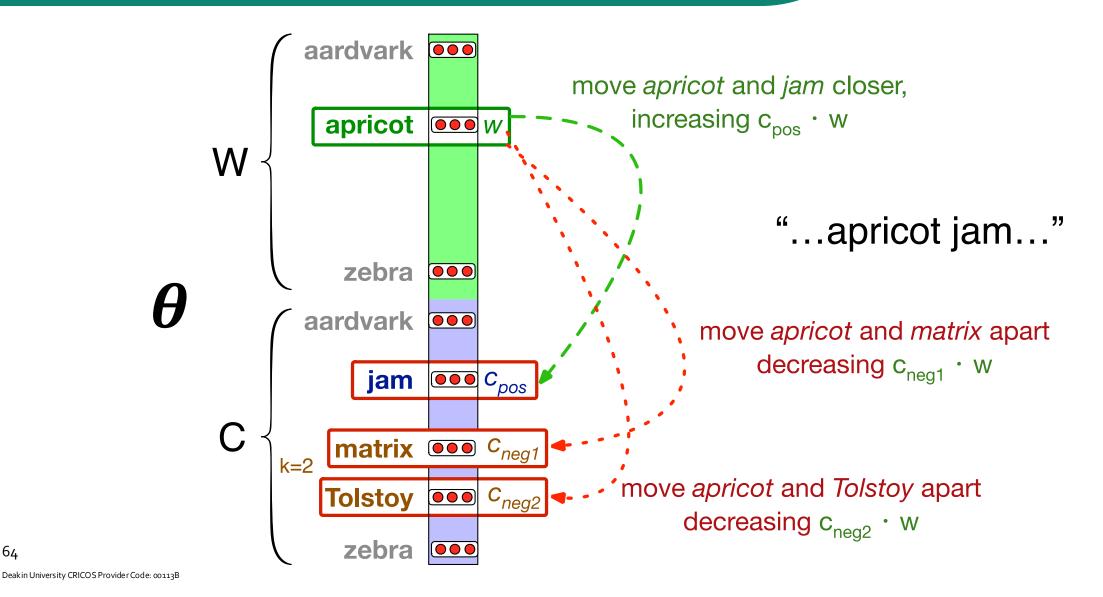


- How to learn?
 - Stochastic gradient descent!

- We'll adjust the word weights to
 - o make the positive pairs more likely
 - o and the negative pairs less likely,
 - over the entire training set.

Intuition of one step of gradient descent





Reminder: gradient descent



- At each step
 - Direction: We move in the reverse direction from the gradient of the loss function
 - Magnitude: we move the value of this gradient $\frac{d}{dw}L(f(x;w),y)$ weighted by a learning rate η
 - Higher learning rate means move w faster

$$w^{t+1} = w^t - h \frac{d}{dw} L(f(x, w), y)$$

The derivatives of the loss function



$$L_{\text{CE}} = -\left[\log \sigma(c_{pos} \cdot w) + \sum_{i=1}^{k} \log \sigma(-c_{neg_i} \cdot w)\right]$$

$$\frac{\partial L_{CE}}{\partial c_{pos}} = [\sigma(c_{pos} \cdot w) - 1]w$$

$$\frac{\partial L_{CE}}{\partial c_{neg}} = [\sigma(c_{neg} \cdot w)]w$$

$$\frac{\partial L_{CE}}{\partial w} = [\sigma(c_{pos} \cdot w) - 1]c_{pos} + \sum_{i=1}^{\kappa} [\sigma(c_{neg_i} \cdot w)]c_{neg_i}$$

Update equation in SGD



Start with randomly initialized C and W matrices, then incrementally do updates

$$c_{pos}^{t+1} = c_{pos}^{t} - \eta [\sigma(c_{pos}^{t} \cdot w^{t}) - 1] w^{t}$$

$$c_{neg}^{t+1} = c_{neg}^{t} - \eta [\sigma(c_{neg}^{t} \cdot w^{t})] w^{t}$$

$$w^{t+1} = w^{t} - \eta \left[[\sigma(c_{pos} \cdot w^{t}) - 1] c_{pos} + \sum_{i=1}^{k} [\sigma(c_{neg_{i}} \cdot w^{t})] c_{neg_{i}} \right]$$

Two sets of embeddings



- SGNS learns two sets of embeddings
 - Target embeddings matrix W
 - Context embedding matrix C
- It's common to just add them together, representing word i as the vector $\mathbf{w_i}$

$$+c_i$$

Summary: How to learn word2vec (skip-gram) embeddings



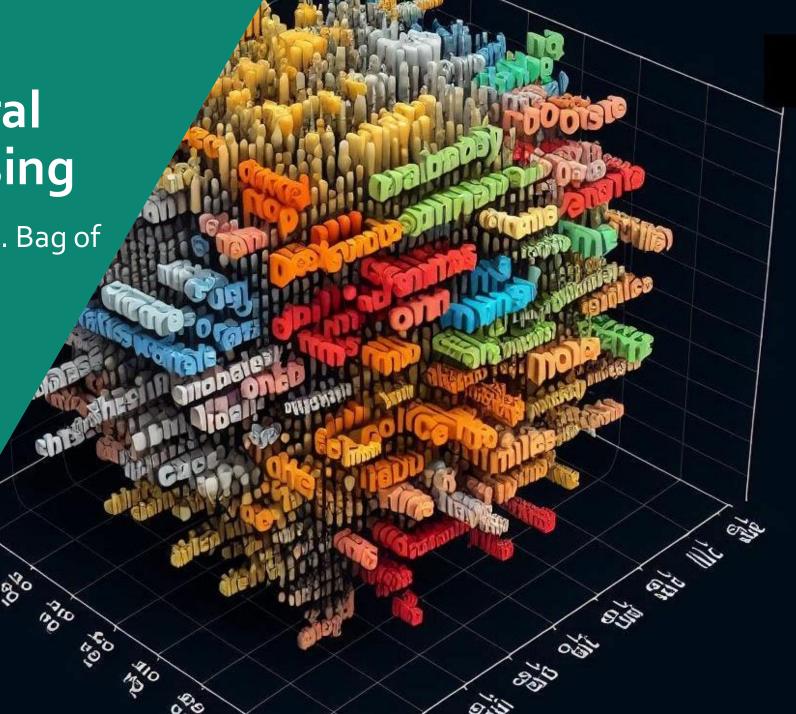
- Start with V random d-dimensional vectors as initial embeddings
- Train a classifier based on embedding similarity
 - Take a corpus and take pairs of words that co-occur as positive examples
 - Take pairs of words that don't co-occur as negative examples
 - o Train the classifier to distinguish these by slowly adjusting all the embeddings to improve the classifier performance
 - Throw away the classifier code and keep the embeddings.



Week 5.6 – Word Embedding vs. Bag of Words

Dr. Mohamed Reda Bouadjenek

School of Information Technology, Faculty of Sci Eng & Built Env



Word Embedding vs. Bag of Words



Traditional Method - Bag of Words Model

Two approaches:

- Either uses one hot encoding.
 - Each word in the vocabulary is represented by one bit position in a HUGE vector.
 - For example, if we have a vocabulary of 10,000 words, and "aardvark" is the 4th word in the dictionary, it would be represented by: [0 0 0 1 0 0 0 0 0].
- Or uses document representation.
 - Each word in the vocabulary is represented by its presence in documents.
 - For example, if we have a corpus of 1M documents, and "Hello" is in 1th, 3th and 5th documents only, it would be represented by: [1 0 1 0 1 0 0 0 0].
- Assumes independence between words.

Word Embeddings

- Stores each word in as a point in space, where it is represented by a dense vector of fixed number of dimensions (generally 300).
 - For example, "Hello" might be represented as: [0.4, -0.11, 0.55, 0.3... 0.1, 0.02].
 - Dimensions are projections along different axes, more of a mathematical concept.
- Unsupervised, built just by reading huge corpus.

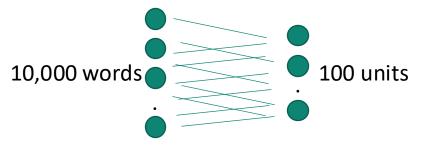
Assumes dependence between words.

Word Embedding vs. Bag of Words



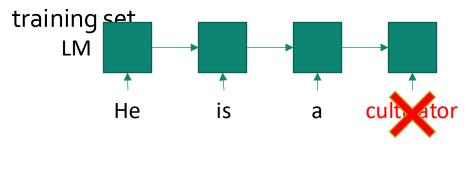
Traditional Method - Bag of Words Model

Requires very large weight matrix for 1st layers.



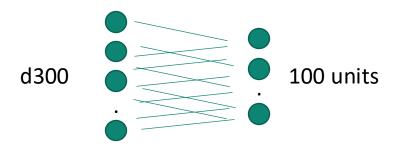
W's size is $10,000 \times 100 = 10^6$

Models not flexible with unseen words in the



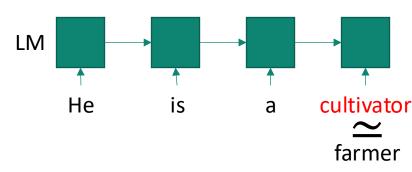
Word Embeddings

A compact weight matrix for 1st layers.



W's size is $300x100 = 3x10^4$

Flexible models with unseen words in the training set.

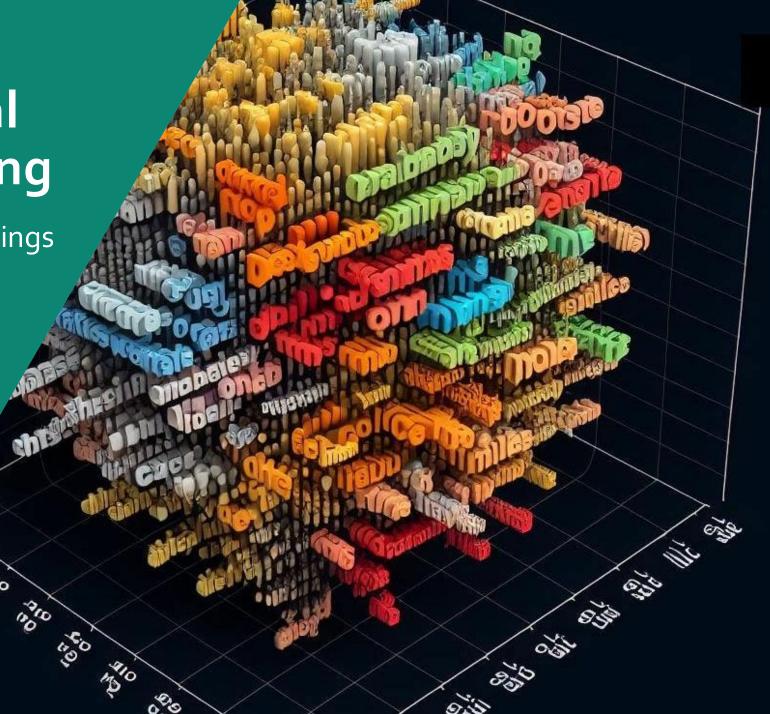




Week 5.7 – Properties of Embeddings

Dr. Mohamed Reda Bouadjenek

School of Information Technology, Faculty of Sci Eng & Built Env



The kinds of neighbors depend on window size



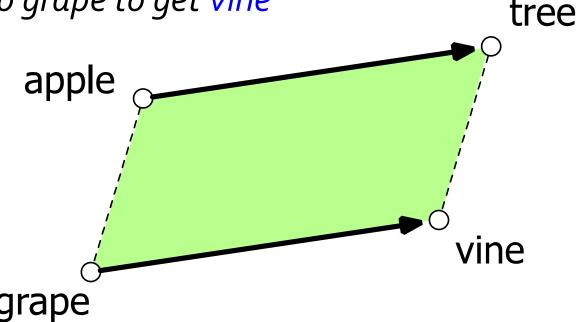
- •Small windows (C= +/- 2) : nearest words are syntactically similar words in same taxonomy
 - Hogwarts nearest neighbors are other fictional schools
 - Sunnydale, Evernight, Blandings
- •Large windows (C = +/-5): nearest words are related words in same semantic field
 - Hogwarts nearest neighbors are Harry Potter world:
 - Dumbledore, half-blood, Malfoy

Analogical relations



- The classic parallelogram model of analogical reasoning (Rumelhart and Abrahamson 1973)
- To solve: "apple is to tree as grape is to _____"

Add tree – apple to grape to get vine

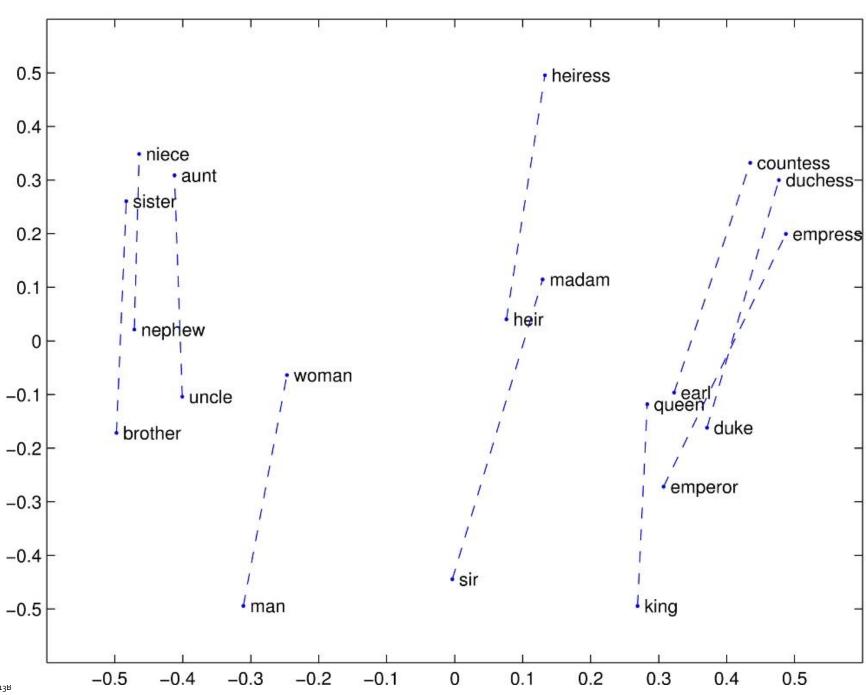


Analogical relations via parallelogram



- The parallelogram method can solve analogies with both sparse and dense embeddings (Turney and Littman 2005, Mikolov et al. 2013b)
- king man + woman is close to queen
- Paris France + Italy is close to Rome
- For a problem a:a*::b:b*, the parallelogram method is:

$$\hat{b}^* = \underset{x}{\operatorname{argmax}} \operatorname{distance}(x, a^* - a + b)$$





Caveats with the parallelogram method



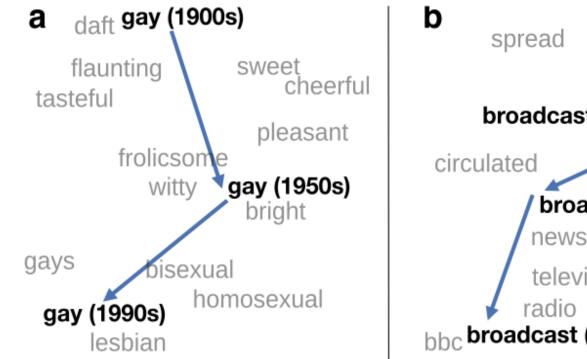
• It only seems to work for frequent words, small distances and certain relations (relating countries to capitals, or parts of speech), but not others. (Linzen 2016, Gladkova et al. 2016, Ethayarajh et al. 2019a)

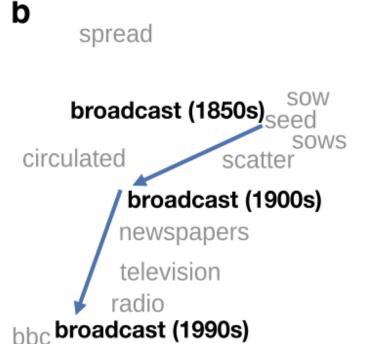
• Understanding analogy is an open area of research (Peterson et al. 2020)

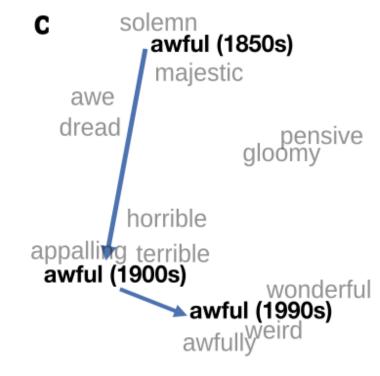
Embeddings as a window onto historical semantics



• Train embeddings on different decades of historical text to see meanings shift ~30 million books, 1850-1990, Google Books data







William L. Hamilton, Jure Leskovec, and Dan Jurafsky. 2016. Diachronic Word Embeddings Reveal Statistical Laws of Semantic Change. Proceedings of ACL.

Embeddings reflect cultural bias!



Bolukbasi, Tolga, Kai-Wei Chang, James Y. Zou, Venkatesh

programmer as woman is to homemaker? debiasing word

Saligrama, and Adam T. Kalai. "Man is to computer

embeddings." In NeurIPS, pp. 4349-4357. 2016.

- Ask "Paris: France:: Tokyo: x"
 - \circ x = Japan
- Ask "father: doctor:: mother: x"
 - o x = nurse
- Ask "man: computer programmer:: woman: x"
 - $\circ x = homemaker$

Algorithms that use embeddings as part of e.g., hiring searches for programmers, might lead to bias in hiring

Historical embedding as a tool to study cultural biases



Garg, N., Schiebinger, L., Jurafsky, D., and Zou, J. (2018). Word embeddings quantify 100 years of gender and ethnic stereotypes. Proceedings of the National Academy of Sciences 115(16), E3635–E3644.

- Compute a gender or ethnic bias for each adjective: e.g., how much closer the
 adjective is to "woman" synonyms than "man" synonyms, or names of particular
 ethnicities
 - Embeddings for **competence** adjective (*smart, wise, brilliant, resourceful, thoughtful, logical*) are biased toward men, a bias slowly decreasing 1960-1990
 - Embeddings for **dehumanizing** adjectives (barbaric, monstrous, bizarre) were biased toward Asians in the 1930s, bias decreasing over the 20th century.
- These match the results of old surveys done in the 1930s

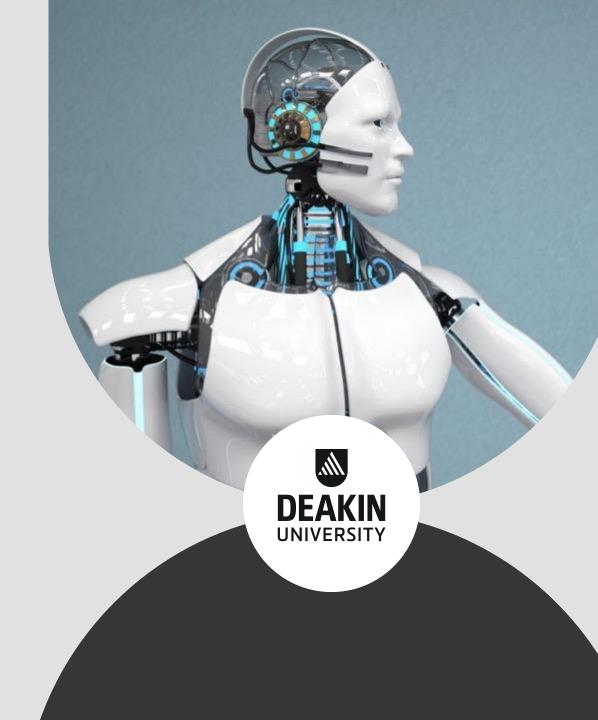
SIT330-770: Natural Language Processing

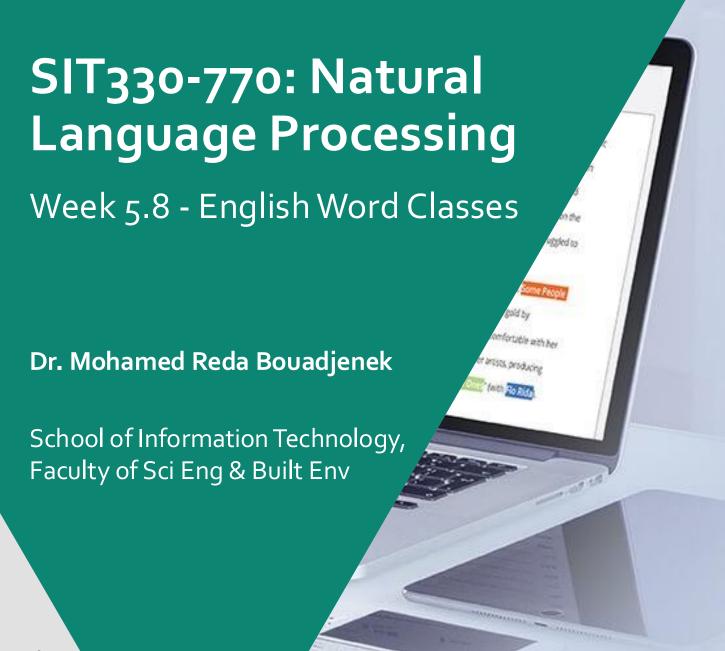
Week 5 - Sequence Labeling

Dr. Mohamed Reda Bouadjenek

School of Information Technology, Faculty of Sci Eng & Built Env

reda.bouadjenek@deakin.edu.au







Parts of Speech



- From the earliest linguistic traditions (Yaska and Panini 5^{th} C. BCE, Aristotle 4^{th} C. BCE), the idea that words can be classified into grammatical categories
 - o part of speech, word classes, POS, POS tags
- 8 parts of speech attributed to Dionysius Thrax of Alexandria (c. 1st C. BCE):
 - o noun, verb, pronoun, preposition, adverb, conjunction, participle, article
 - These categories are relevant for NLP today.

Two classes of words: Open vs. Closed



Closed class words

- Relatively fixed membership
- Usually function words: short, frequent words with grammatical function
 - determiners: a, an, the
 - pronouns: she, he, I
 - prepositions: on, under, over, near, by, ...

Open class words

- Usually content words: Nouns, Verbs, Adjectives, Adverbs
 - Plus interjections: oh, ouch, uh-huh, yes, hello
- New nouns and verbs like *iPhone* or *to fax*

Open class ("content") words

Nouns

Proper

Janet Italy Common

cat, cats mango

Verbs

Main

eat went Adjectives old green tasty

Adverbs slowly yesterday

Numbers

122,312

one

Interjections Ow hello

... more

Closed class ("function")

Determiners the some

Conjunctions and or

Pronouns they its

Auxiliary

can had Prepositions to with

Particles off up

... more

Part-of-Speech Tagging



- Assigning a part-of-speech to each word in a text.
- Words often have more than one POS.
- book:
 - VERB: (Book that flight)
 - NOUN: (Hand me that **book**).

"Universal Dependencies" Tagset

	Tag	Description	Example
	ADJ	Adjective: noun modifiers describing properties	red, young, awesome
Class	ADV	Adverb: verb modifiers of time, place, manner	very, slowly, home, yesterday
CI	NOUN	words for persons, places, things, etc.	algorithm, cat, mango, beauty
Open	VERB	words for actions and processes	draw, provide, go
O	PROPN	Proper noun: name of a person, organization, place, etc	Regina, IBM, Colorado
	INTJ	Interjection: exclamation, greeting, yes/no response, etc.	oh, um, yes, hello
	ADP	Adposition (Preposition/Postposition): marks a noun's	in, on, by under
<u>s</u>		spacial, temporal, or other relation	
Closed Class Words	AUX	Auxiliary: helping verb marking tense, aspect, mood, etc.,	can, may, should, are
	CCONJ	Coordinating Conjunction: joins two phrases/clauses	and, or, but
ass	DET	Determiner: marks noun phrase properties	a, an, the, this
[]	NUM	Numeral	one, two, first, second
sed	PART	Particle: a preposition-like form used together with a verb	up, down, on, off, in, out, at, by
$\Box 10$	PRON	Pronoun: a shorthand for referring to an entity or event	she, who, I, others
	SCONJ	Subordinating Conjunction: joins a main clause with a	that, which
		subordinate clause such as a sentential complement	
er	PUNCT	Punctuation	; ,()
Other	SYM	Symbols like \$ or emoji	\$, %
	X	Other	asdf, qwfg

88

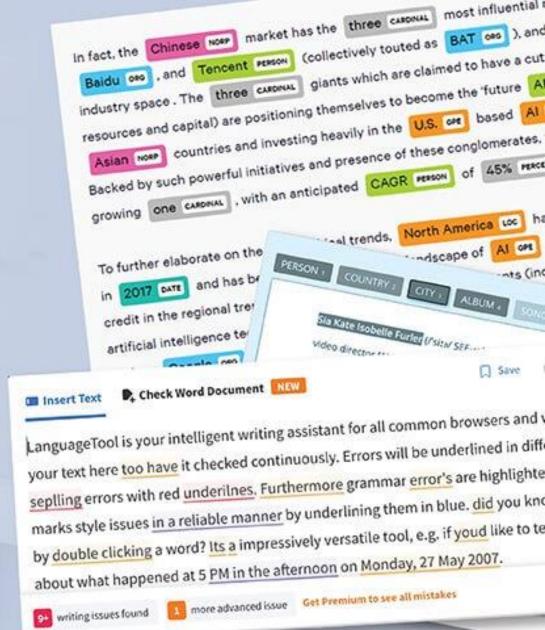
Sample "Tagged" English sentences



There/PRO were/VERB 70/NUM children/NOUN there/ADV ./PUNC

 Preliminary/ADJ findings/NOUN were/AUX reported/VERB in/ADP today/NOUN 's/PART New/PROPN England/PROPN Journal/PROPN of/ADP Medicine/PROPN

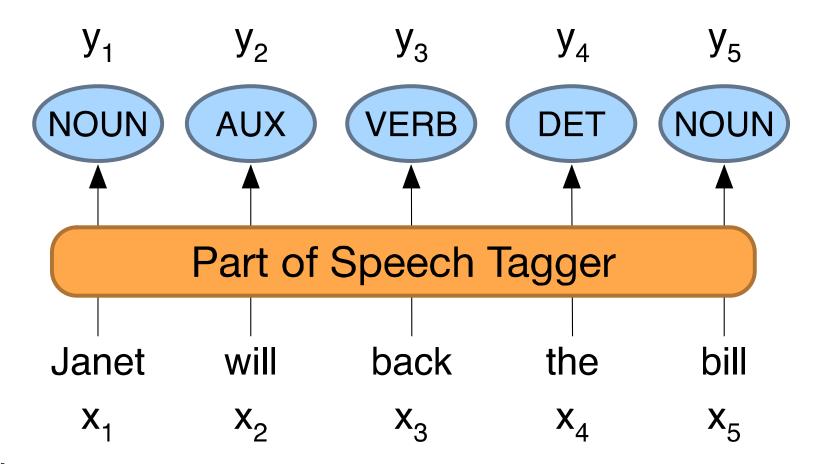




Part-of-Speech Tagging



• Map from sequence $x_1, ..., x_n$ of words to $y_1, ..., y_n$ of POS tags



Why Part of Speech Tagging?



- Can be useful for other NLP tasks
 - Parsing: POS tagging can improve syntactic parsing
 - MT: reordering of adjectives and nouns (say from Spanish to English)
 - Sentiment or affective tasks: may want to distinguish adjectives or other POS
 - Text-to-speech (how do we pronounce "lead" or "object"?)
- Or linguistic or language-analytic computational tasks
 - Need to control for POS when studying linguistic change like creation of new words, or meaning shift
 - Or control for POS in measuring meaning similarity or difference

How difficult is POS tagging in English?



- Roughly 15% of word types are ambiguous
 - Hence 85% of word types are unambiguous
 - Janet is always PROPN, hesitantly is always ADV
- But those 15% tend to be very common.
- So ~60% of word tokens are ambiguous
- E.g., back
 - earnings growth took a back/ADJ seat
 - o a small building in the back/NOUN
 - a clear majority of senators back/VERB the bill
 - enable the country to buy back/PART debt
 - I was twenty-one back/ADV then

POS tagging performance in English



- How many tags are correct? (Tag accuracy)
 - About 97%
 - Hasn't changed in the last 10+ years
 - HMMs, CRFs, BERT perform similarly.
 - Human accuracy about the same
- But baseline is 92%!
 - Baseline is performance of stupidest possible method
 - "Most frequent class baseline" is an important baseline for many tasks
 - Tag every word with its most frequent tag
 - (and tag unknown words as nouns)
 - Partly easy because
 - Many words are unambiguous

Sources of information for POS tagging



Janet will back the bill

AUX/NOUN/VERB? NOUN/VERB?

- Prior probabilities of word/tag
 - "will" is usually an AUX
- Identity of neighboring words
 - o "the" means the next word is probably not a verb
- Morphology and wordshape:

 \circ Prefixes unable: un- \rightarrow ADJ

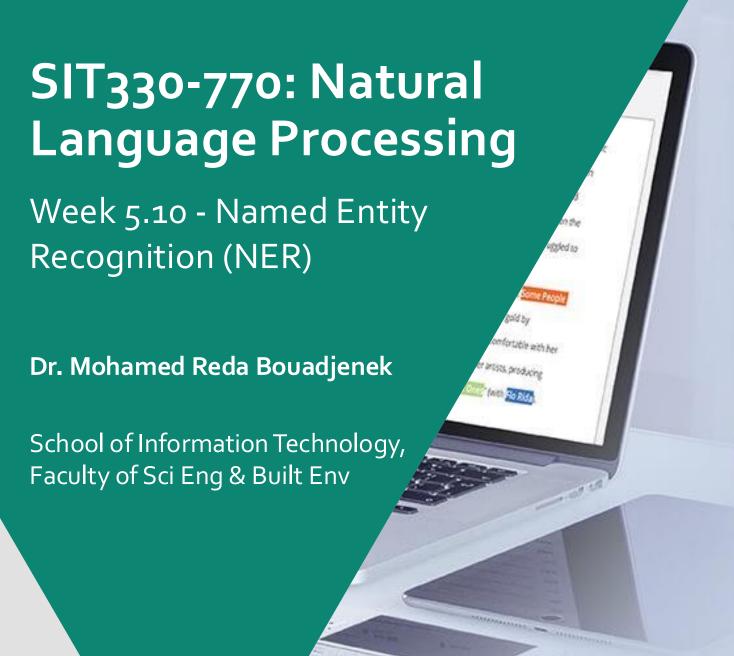
 \circ Suffixes importantly: -ly \rightarrow ADV

 \circ Capitalization Janet: CAP \rightarrow PROPN

Standard algorithms for POS tagging



- Supervised Machine Learning Algorithms:
- Hidden Markov Models
- Conditional Random Fields (CRF)/ Maximum Entropy Markov Models (MEMM)
- Neural sequence models (RNNs or Transformers)
- Large Language Models (like BERT), finetuned
- All required a hand-labeled training set, all about equal performance (97% on English)
- All make use of information sources we discussed
- Via human created features: HMMs and CRFs
- Via representation learning: Neural LMs





Named Entities



- Named entity, in its core usage, means anything that can be referred to with a proper name. Most common 4 tags:
 - PER (Person): "Marie Curie"
 - LOC (Location): "New York City"
 - ORG (Organization): "Stanford University"
 - GPE (Geo-Political Entity): "Boulder, Colorado"
- Often multi-word phrases
- But the term is also extended to things that aren't entities:
 - o dates, times, prices

Named Entity tagging



- The task of named entity recognition (NER):
 - o find spans of text that constitute proper names
 - o tag the type of the entity.

NER output



Citing high fuel prices, [ORG United Airlines] said [TIME Friday] it has increased fares by [MONEY \$6] per round trip on flights to some cities also served by lower-cost carriers. [ORG American Airlines], a unit of [ORG AMR Corp.], immediately matched the move, spokesman [PER Tim Wagner] said. [ORG United], a unit of [ORG UAL Corp.], said the increase took effect [TIME Thursday] and applies to most routes where it competes against discount carriers, such as [LOC Chicago] to [LOC Dallas] and [LOC Denver] to [LOC San Francisco].

Why NER?



- Sentiment analysis: consumer's sentiment toward a particular company or person?
- Question Answering: answer questions about an entity?
- Information Extraction: Extracting facts about entities from text.

Why NER is hard



Segmentation

- In POS tagging, no segmentation problem since each word gets one tag.
- In NER we have to find and segment the entities!

Type ambiguity

[PER Washington] was born into slavery on the farm of James Burroughs. [ORG Washington] went up 2 games to 1 in the four-game series. Blair arrived in [LOC Washington] for what may well be his last state visit. In June, [GPE Washington] passed a primary seatbelt law.

BIO Tagging



 How can we turn this structured problem into a sequence problem like POS tagging, with one label per word?

• [PER Jane Villanueva] of [ORG United], a unit of [ORG United Airlines Holding], said the fare applies to the [LOC Chicago] route.

BIO Tagging



[PER Jane Villanueva] of [ORG United], a unit of [ORG United Airlines

Holding], said the fare applies to the [LOC Chicago] route.

Words	BIO Label
Jane	B-PER
Villanueva	I-PER
of	O
United	B-ORG
Airlines	I-ORG
Holding	I-ORG
discussed	O
the	O
Chicago	B-LOC
route	O
•	O

Now we have one tag per token!!!

BIO Tagging



B: token that *begins* a span

I: tokens *inside* a span

O: tokens outside of any span

of tags (where n is #entity types):

10 tag,

n B tags,

n I tags

total of 2n+1

Words	BIO Label	
Jane	B-PER	
Villanueva	I-PER	
of	O	
United	B-ORG	
Airlines	I-ORG	
Holding	I-ORG	
discussed	O	
the	O	
Chicago	B-LOC	
route	O	
•	O	

BIO Tagging variants: IO and BIOES



• [PER Jane Villanueva] of [ORG United], a unit of [ORG United Airlines

Holding], said the fare applies to the [LOC Chicago] route.

Words	IO Label	BIO Label	BIOES Label
Jane	I-PER	B-PER	B-PER
Villanueva	I-PER	I-PER	E-PER
of	O	O	O
United	I-ORG	B-ORG	B-ORG
Airlines	I-ORG	I-ORG	I-ORG
Holding	I-ORG	I-ORG	E-ORG
discussed	O	O	O
the	O	O	O
Chicago	I-LOC	B-LOC	S-LOC
route	O	O	O
osp. •	O	O	O

Standard algorithms for NER



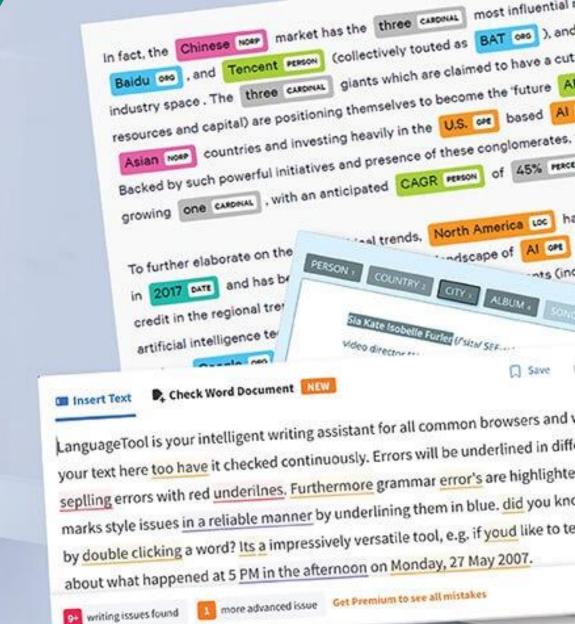
- Supervised Machine Learning given a human-labeled training set of text annotated with tags
 - Hidden Markov Models
 - Conditional Random Fields (CRF)/ Maximum Entropy Markov Models (MEMM)
 - Neural sequence models (RNNs or Transformers)
 - Large Language Models (like BERT), finetuned



Week 5.11 – Hidden Markov Model (HMM) Part-of-Speech Tagging

Dr. Mohamed Reda Bouadjenek

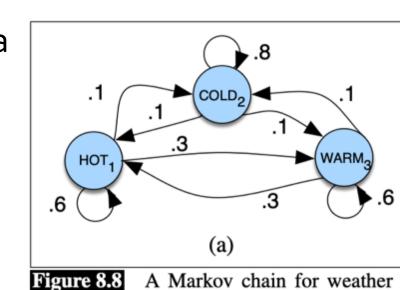
School of Information Technology, Faculty of Sci Eng & Built Env



Introduction to Markov Chains



- A Markov chain models the probabilities of state sequences, each drawn from a specific set.
- It assumes the future state depends only on the current state, not any prior ones.
- Markov chains are used to predict various phenomena
 - E.g., modeling weather patterns or word sequences.



Markov Chain Representation



$$Q = q_1q_2 \dots q_N$$
 a set of N states $A = a_{11}a_{12} \dots a_{N1} \dots a_{NN}$ a transition probability matrix A , each a_{ij} representing the probability of moving from state i to state j , s.t.
$$\sum_{j=1}^{n} a_{ij} = 1 \quad \forall i$$
 an initial probability distribution over states. π_i is the probability that the Markov chain will start in state i . Some states j may have $\pi_j = 0$, meaning that they cannot be initial states. Also, $\sum_{i=1}^{n} \pi_i = 1$

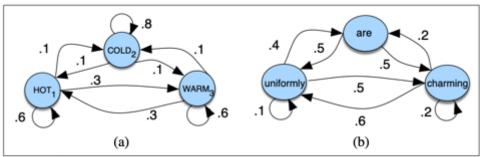


Figure 8.8 A Markov chain for weather (a) and one for words (b), showing states and transitions. A start distribution π is required; setting $\pi = [0.1, 0.7, 0.2]$ for (a) would mean a probability 0.7 of starting in state 2 (cold), probability 0.1 of starting in state 1 (hot), etc.

Markov Assumption:

o Formally stated as: $P(q_i=a|q_1...q_{i-1})=P(q_i=a|q_{i-1})$ implying that when predicting the future, only the present state matters

The Hidden Markov Model



- A Markov chain computes probabilities for sequences of observable events.
- But often, the events of interest are hidden.
 - **Example:** Part-of-speech tags in text—hidden because we don't observe them directly.
- **Solution:** Hidden Markov Model (HMM) handles both observed and hidden events.
 - HMMs augment Markov chains

Probabilistic Sequence Modeling with HMMs



- A Hidden Markov Models (HMM) is a probabilistic sequence model that, given a sequence of units (words, letters, morphemes, sentences, etc.), computes a probability distribution over possible sequences of labels.
 - HMMs determine the likelihood of different label sequences and select the most probable sequence based on the observed data.
 - HMM is based on augmenting the Markov chain

Input and Assumptions



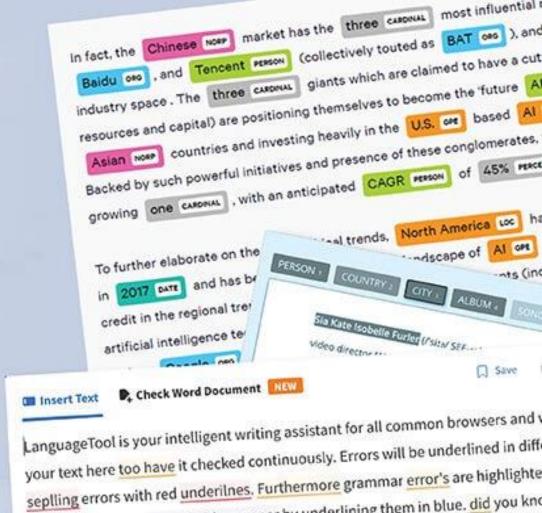
- Input (O): Sequence of observations (o_1 , o_2 , ..., o_T) drawn from vocabulary V.
- Assumptions of first-order HMM:
 - Markov Assumption:
 - \circ Probability of state q_i depends only on the previous state (q_{i-1}) .
 - $P(q_i|q_1...q_{i-1}) = P(q_i|q_{i-1})$
 - Output Independence:
 - Probability of observation o_i depends only on the state that produced it q_i
 - $P(o_i | q_1, ..., q_T, o_1, ..., o_i, ..., o_T) = P(o_i | q_i)$



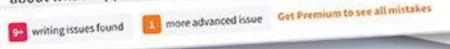
Week 5.12 — The components of an HMM tagger

Dr. Mohamed Reda Bouadjenek

School of Information Technology, Faculty of Sci Eng & Built Env



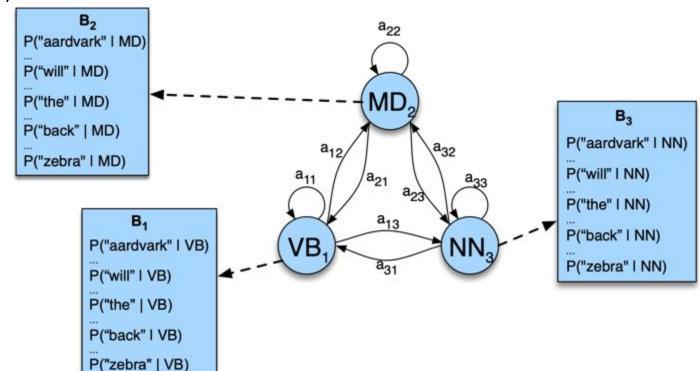
marks style issues in a reliable manner by underlining them in blue. did you kno by double clicking a word? Its a impressively versatile tool, e.g. if youd like to te about what happened at 5 PM in the afternoon on Monday, 27 May 2007.



Components of an HMM Tagger



- An HMM tagger consists of two main components:
 - Matrix A which represents the tag transition.
 - Matrix B which represents emission probabilities.



The A Matrix - Transition Probabilities



- The A matrix encapsulates the tag transition probabilities, $P(t_i|t_{i-1})$, which express how likely a tag follows its predecessor.
 - Example:
 - The modal verb "will" commonly precedes the base form of a verb (VB), as in "will race", leading to a high transition probability.
 - o These probabilities are derived using maximum MLE by counting tag occurrences in a labeled corpus.
- Calculating Transition Probabilities:
 - In the WSJ corpus example, the modal verb tag (MD) is observed 13,124 times.
 - Out of these, MD transitions to a base verb (VB) 10,471 times.
 - Using MLE, we estimate $P(VB|MD) = C(MD, VB) / C(MD) = 10,471 / 13,124 \approx 0.80$.

The B Matrix - Emission Probabilities



- The B matrix contains emission probabilities, $P(w_i|t_i)$, which quantify the likelihood of a word being tagged with a specific tag.
- Emission Probability Calculation
 - To calculate emission probabilities, we count how often a word occurs with a particular tag in a corpus.
 - For instance, the MD tag associated with the word 'will' occurs 4,046 times in the WSJ corpus.
 - Hence, P(will|MD) is calculated as C(MD, will) / C(MD) = 4,046 / $13,124 \approx 0.31$.

Components of HMM



$$Q = q_1 q_2 \dots q_N$$
 a set of N states

$$A = a_{11} \dots a_{ij} \dots a_{NN}$$

$$B = b_i(o_t)$$

$$\pi = \pi_1, \pi_2, ..., \pi_N$$

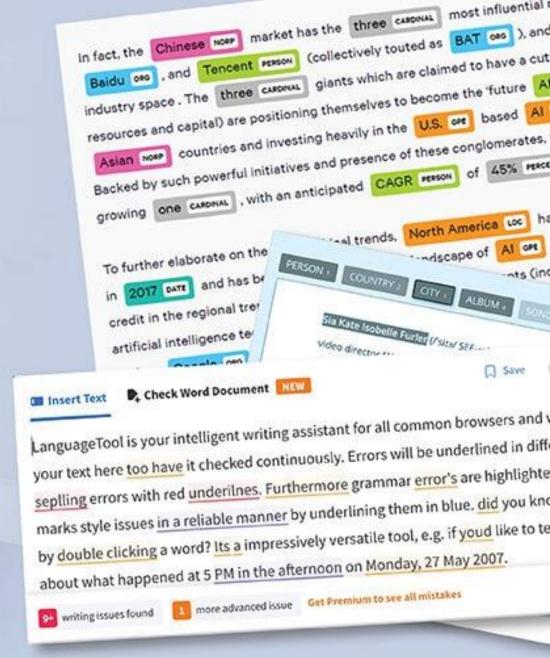
a transition probability matrix A, each a_{ij} representing the probability of moving from state i to state j, s.t. $\sum_{j=1}^{N} a_{ij} = 1 \quad \forall i$ a sequence of **observation likelihoods**, also called **emission probabilities**, each expressing the probability of an observation o_t (drawn from a

an initial probability distribution over states. π_i is the probability that the Markov chain will start in state *i*. Some states *j* may have $\pi_i = 0$,

meaning that they cannot be initial states. Also, $\sum_{i=1}^{n} \pi_i = 1$

vocabulary $V = v_1, v_2, ..., v_V$) being generated from a state q_i





Decoding with Hidden Markov Models



- Decoding is the process of determining the most probable sequence of hidden states (tags) based on observed data.
 - Given a sequence of observations $O = o_1$, o_2 , ..., o_T , decoding aims to find the most probable sequence of states $Q = q_1 q_2 \dots q_T$.
 - The input is an HMM λ = (A , B), with **A** being the transition probabilities and **B** the emission probabilities.

$$\hat{t}_{1:n} = \underset{t_1...t_n}{\operatorname{argmax}} P(t_1...t_n | w_1...w_n)$$

Decoding with Hidden Markov Models (i)



$$\hat{t}_{1:n} = \mathop{\mathrm{argmax}}_{t_1...t_n} P(t_1...t_n|w_1...w_n)$$
 MAP is "maximum a posteriori" = most likely sequence

$$\hat{t}_{1:n} = \operatorname*{argmax}_{t_1 \dots t_n} rac{P(w_1 \dots w_n | t_1 \dots t_n) P(t_1 \dots t_n)}{P(w_1 \dots w_n)}$$
 Bayes Rule

$$\hat{t}_{1:n} = \underset{t_1...t_n}{\operatorname{argmax}} P(w_1 ... w_n | t_1 ... t_n) P(t_1 ... t_n)$$

Dropping the denominator

Decoding with Hidden Markov Models (ii)



$$\hat{t}_{1:n} = \operatorname*{argmax}_{t_1...t_n} P(w_1...w_n|t_1...t_n) P(t_1...t_n)$$

- HMM taggers make two further simplifying assumptions.
 - The probability of a word appearing depends only on its own tag and is independent of neighboring words and tags: $P(w_1 \dots w_n | t_1 \dots t_n) \approx \prod_{i=1}^n P(w_i | t_i)$
 - The second assumption, the bigram assumption, is that the probability of a tag is dependent only on the previous tag, rather than the entire tag sequence;

$$P(t_1...t_n) \approx \prod_{i=1}^n P(t_i|t_{i-1})$$

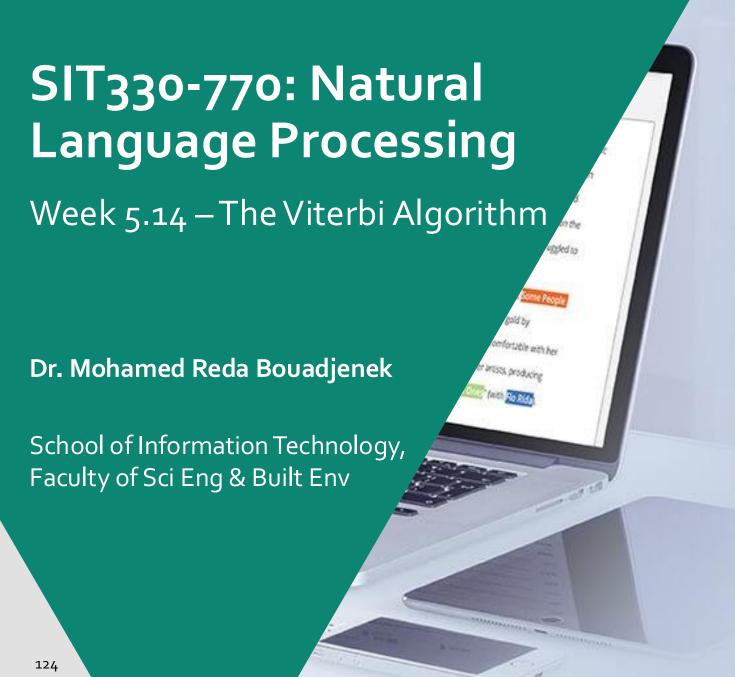
Decoding with Hidden Markov Models (iii)

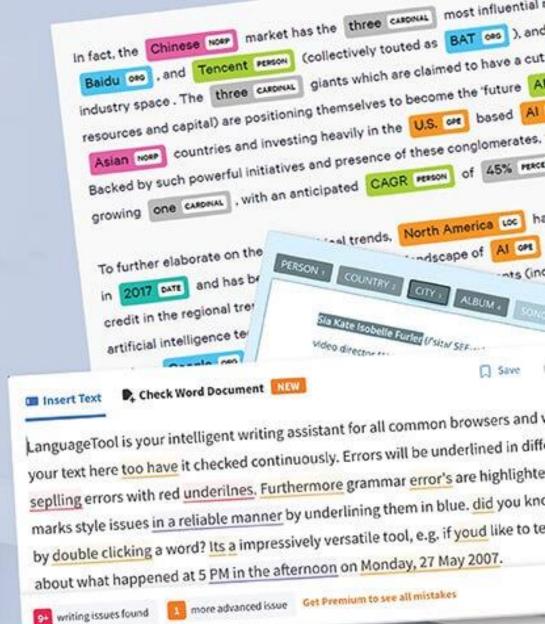


• Plugging the simplifying assumptions results in the following equation for the most probable tag sequence from a bigram tagger:

$$\hat{t}_{1:n} = \underset{t_1...t_n}{\operatorname{argmax}} P(t_1...t_n|w_1...w_n) \approx \underset{t_1...t_n}{\operatorname{argmax}} \prod_{i=1}^n \underbrace{P(w_i|t_i)}_{P(t_i|t_{i-1})}$$

• The two parts correspond neatly to the **B** emission probability and **A** transition probability that we defined previously!





Computing the most probable sequence of tags



- A brute force approach to identify the most probable sequence of tags faces exponential complexity
 - This method is impractical for large datasets or real-time applications.
- Solution: The Viterbi algorithm 1967
 - Leverages dynamic programming, streamlining the process by breaking the problem into manageable sub-problems
 - This approach significantly reduces computational demands and enhances processing speed, making it viable for complex tasks in real-world scenarios



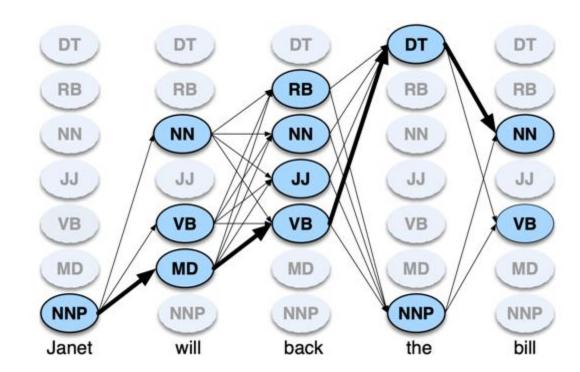
Andrew Viterbi

The Viterbi Algorithm (i)



- The decoding algorithm for HMMs is the Viterbi algorithm
 - o As an instance of dynamic programming, Viterbi resembles the dynamic programming minimum edit distance algorithm
- The Viterbi algorithm first sets up a probability matrix or lattice:
 - Columns as observables (words of a sentence in the same sequence as in sentence)
 - Rows as hidden states (all possible POS Tags are known)

tag the sentence
Janet will back the bill



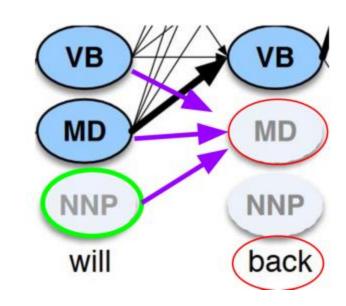
The Viterbi Algorithm (ii)



- Each cell of the matrix is represented by $V_t(j)$ (Viterbi value for t: column, j: row) having the probability that the HMM is in **state** j (present POS Tag) after seeing the **first** t **observations** (past words for which matrix (cell) values has been calculated) and passing through the most **probable state sequence** (previous POS Tag) q_1, \dots, q_{t-1}
- Computed by recursively taking the most probable path that could lead us to this cell

$$v_t(j) = \max_{i=1}^{N} v_{t-1}(i) a_{ij} b_j(o_t)$$

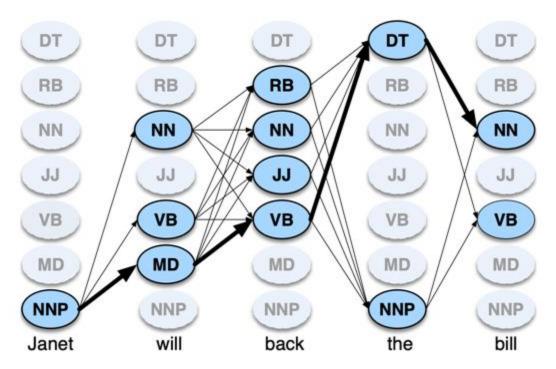
 $v_{t-1}(i)$ the **previous Viterbi path probability** from the previous time step the **transition probability** from previous state q_i to current state q_j the **state observation likelihood** of the observation symbol o_t given the current state j



The Viterbi Algorithm (iii)



• Each cell of the matrix is represented by $V_t(j)$ (Viterbi value for t: column, j: row) having the probability that the HMM is in **state** j (present POS Tag) after seeing the **first** t **observations** (past words for which matrix (cell) values has been calculated) and passing through the most **probable state sequence** (previous POS Tag) q_1 q_{t-1}



A sketch of the matrix for Janet will back the bill, showing the possible tags (q_i) for each word and highlighting the path corresponding to the correct tag sequence through the hidden states

States (parts of speech) which have a zero probability of generating a particular word according to the B matrix (such as the probability that a determiner DT will be realized as Janet) are greyed out

Working Example (i)



Janet will back the bill → Janet/NNP will/MD back/VB the/DT bill/NN

The A transition probabilities $P(t_i|t_{i-1})$ computed from the WSJ corpus without smoothing

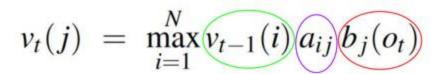
	NNP	MD	VB	JJ	NN	RB	DT
<s></s>	0.2767	0.0006	0.0031	0.0453	0.0449	0.0510	0.2026
NNP	0.3777	0.0110	0.0009	0.0084	0.0584	0.0090	0.0025
MD	0.0008	0.0002	0.7968	0.0005	0.0008	0.1698	0.0041
VB	0.0322	0.0005	0.0050	0.0837	0.0615	0.0514	0.2231
JJ	0.0366	0.0004	0.0001	0.0733	0.4509	0.0036	0.0036
NN	0.0096	0.0176	0.0014	0.0086	0.1216	0.0177	0.0068
RB	0.0068	0.0102	0.1011	0.1012	0.0120	0.0728	0.0479
DT	0.1147	0.0021	0.0002	0.2157	0.4744	0.0102	0.0017

Observation likelihoods B computed from the WSJ corpus without smoothing, simplified slightly

	Janet	will	back	the	bill
NNP	0.000032	0	0	0.000048	0
MD	0	0.308431	0	0	0
VB	0	0.000028	0.000672	0	0.000028
JJ	0	0	0.000340	0	0
NN	0	0.000200	0.000223	0	0.002337
RB	0	0	0.010446	0	0
DT	0	0	0	0.506099	0

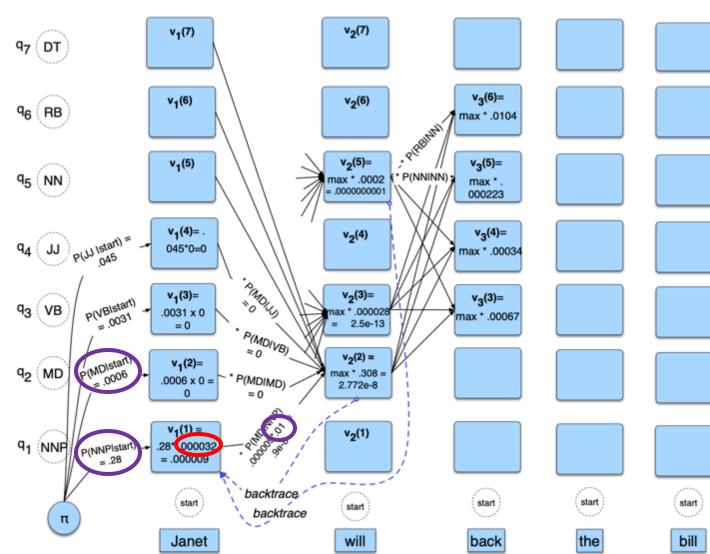
Working Example (ii)





	NND	MD	VB	JJ	NN	RB	DT
<s></s>	0.2767	0.0006	0.0031	0.0453	0.0449	0.0510	0.2026
NNP	0.3777	0.0110	0.0009	0.0084	0.0584	0.0090	0.0025
MD	0.0008	0.0002	0.7968	0.0005	0.0008	0.1698	0.0041
VB	0.0322	0.0005	0.0050	0.0837	0.0615	0.0514	0.2231
JJ	0.0366	0.0004	0.0001	0.0733	0.4509	0.0036	0.0036
NN	0.0096	0.0176	0.0014	0.0086	0.1216	0.0177	0.0068
RB	0.0068	0.0102	0.1011	0.1012	0.0120	0.0728	0.0479
DT	0.1147	0.0021	0.0002	0.2157	0.4744	0.0102	0.0017

	Janet	will	back	the	bill
NNP	0.00003	2 0	0	0.000048	0
MD	0	0.308431	0	0	0
VB	0	0.000028	0.000672	0	0.000028
JJ	0	0	0.000340	0	0
NN	0	0.000200	0.000223	0	0.002337
RB	0	0	0.010446	0	0
DT	0	0	0	0.506099	0



Evaluation Metrics for Named Entity Recognition (NER)



NER Evaluation Basics:

- Unlike POS tagging, evaluated on accuracy, NER uses recall, precision, and F1 score.
- Recall measures correctly identified entities against all actual entities.
- Precision counts correct labels against all labeling attempts.
- The F1 score provides a balance between precision and recall, serving as a single metric for accuracy.

Challenges in NER:

- NER systems treat entities as single units for evaluation, leading to challenges not seen in POS tagging.
- The system's ability to correctly identify entire entities, such as 'Jane Villanueva', impacts evaluation outcomes.
- Mismatches in entity recognition across training and test data can skew results.